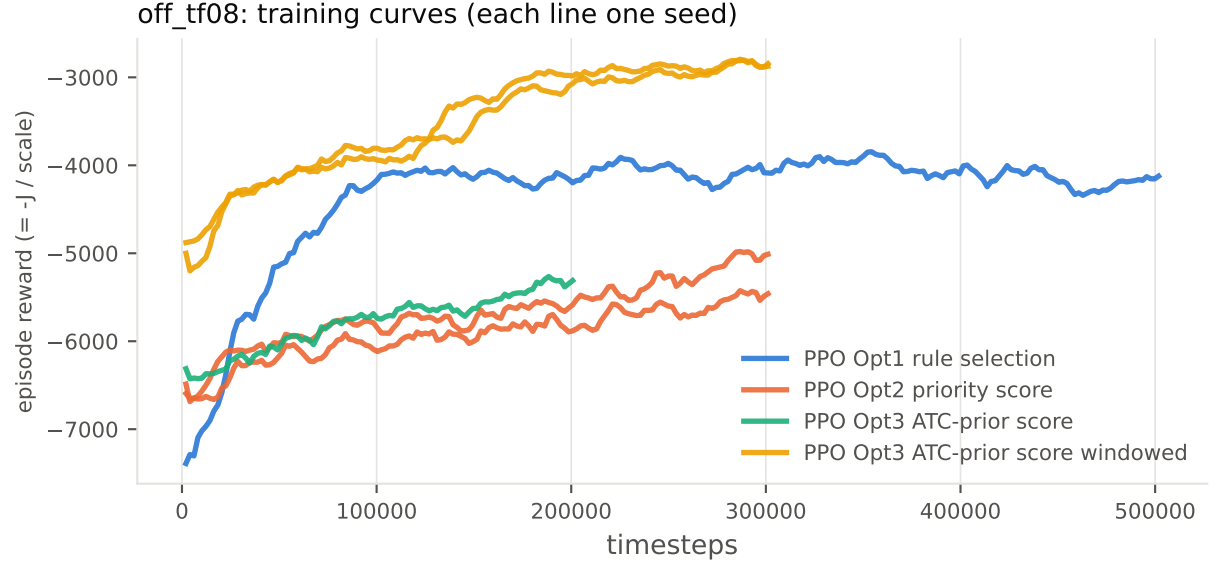
off_tf08: training curves (each line one seed)
episode reward (= -J / scale)
−3000
−4000
−5000
−6000
−7000
0
100000
200000
300000
400000
500000
timesteps
PPO Opt1 rule selection
PPO Opt2 priority score
PPO Opt3 ATC-prior score
PPO Opt3 ATC-prior score windowed