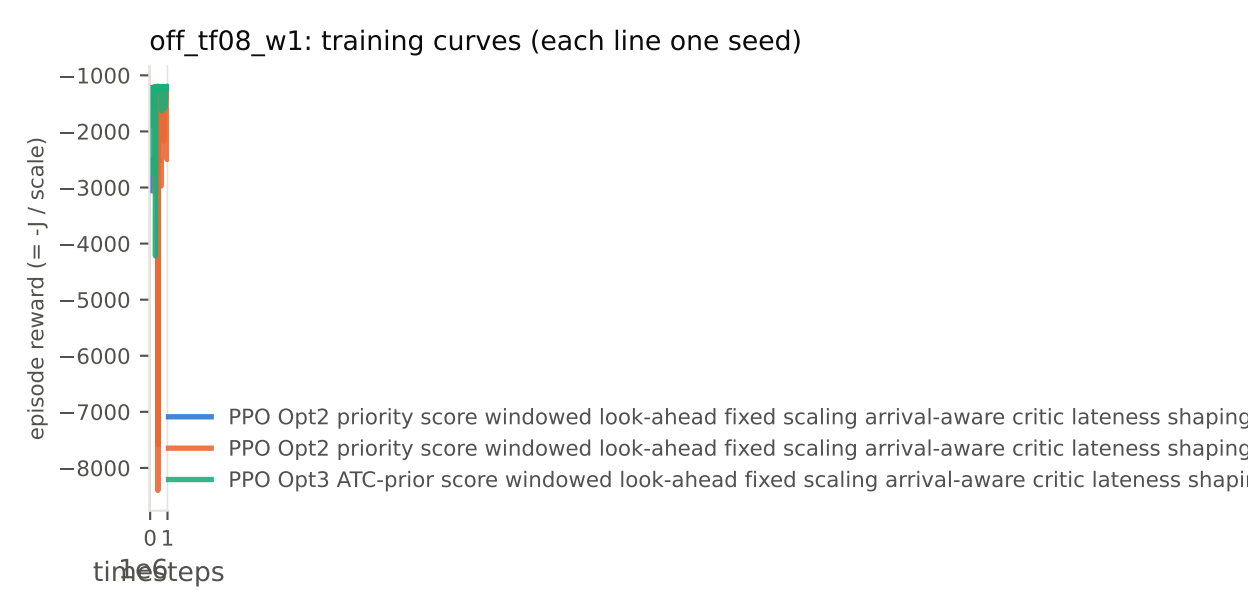
off_tf08_w1: training curves (each line one seed)
episode reward (= -J / scale)
−1000
−2000
−3000
−4000
−5000
−6000
−7000
−8000
0
1
1e6
timesteps
PPO Opt2 priority score windowed look-ahead fixed scaling arrival-aware critic lateness shaping
PPO Opt2 priority score windowed look-ahead fixed scaling arrival-aware critic lateness shaping
PPO Opt3 ATC-prior score windowed look-ahead fixed scaling arrival-aware critic lateness shapi