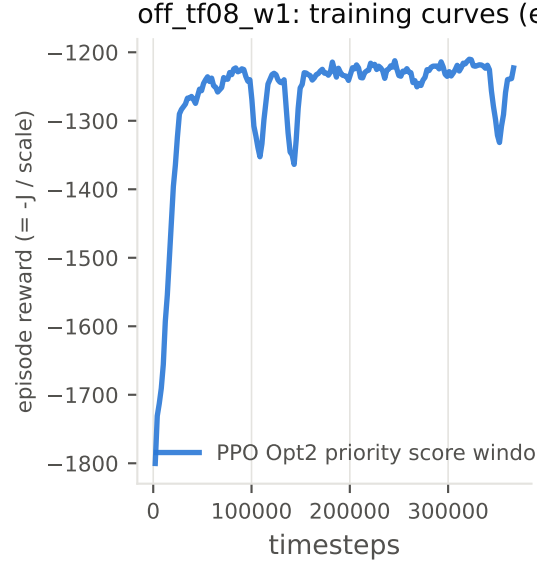
off_tf08_w1: training curves (each line one seed)
episode reward (= -J / scale)
−1200
−1300
−1400
−1500
−1600
−1700
−1800
0
100000
200000
300000
timesteps
PPO Opt2 priority score windowed look-ahead fixed scaling arrival-aware critic lateness shaping