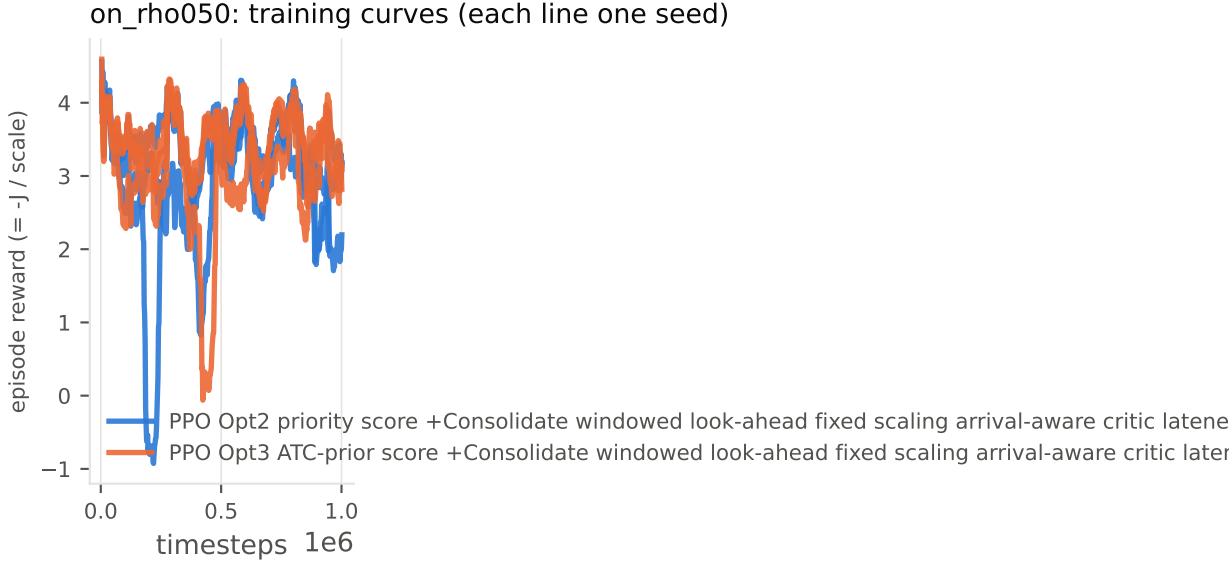
on_rho050: training curves (each line one seed)
episode reward (= -J / scale)
4
3
2
1
0
−1
0.0
0.5
1.0
timesteps
1e6
PPO Opt2 priority score +Consolidate windowed look-ahead fixed scaling arrival-aware critic latene
PPO Opt3 ATC-prior score +Consolidate windowed look-ahead fixed scaling arrival-aware critic later