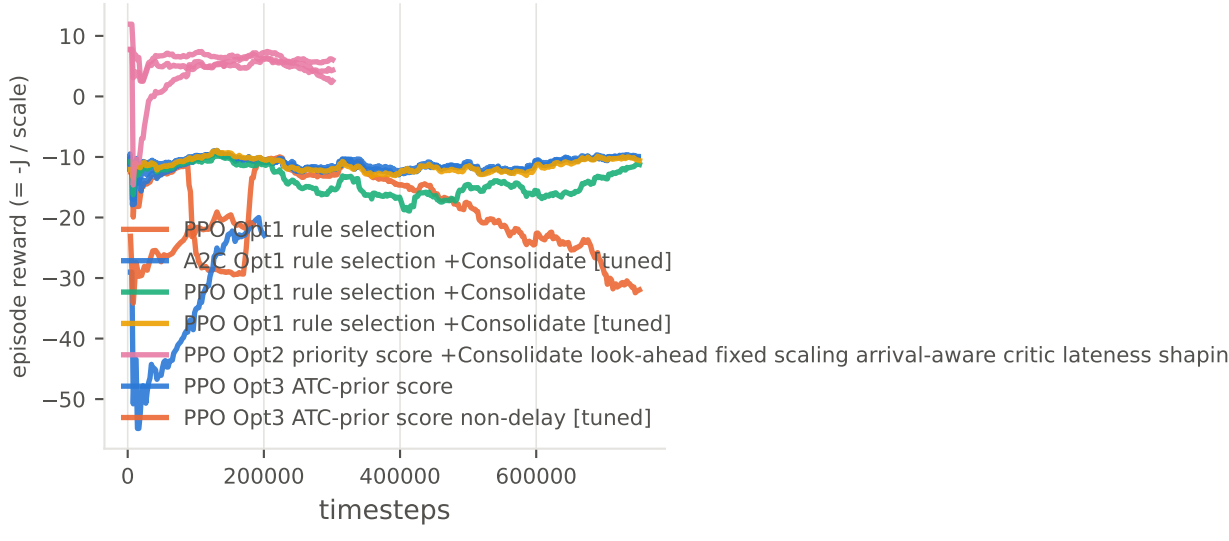
on_rho075: training curves (each line one seed)
episode reward (= -J / scale)
10
0
−10
−20
−30
−40
−50
0
200000
400000
600000
timesteps
PPO Opt1 rule selection
A2C Opt1 rule selection +Consolidate [tuned]
PPO Opt1 rule selection +Consolidate
PPO Opt1 rule selection +Consolidate [tuned]
PPO Opt2 priority score +Consolidate look-ahead fixed scaling arrival-aware critic lateness shapin
PPO Opt3 ATC-prior score
PPO Opt3 ATC-prior score non-delay [tuned]