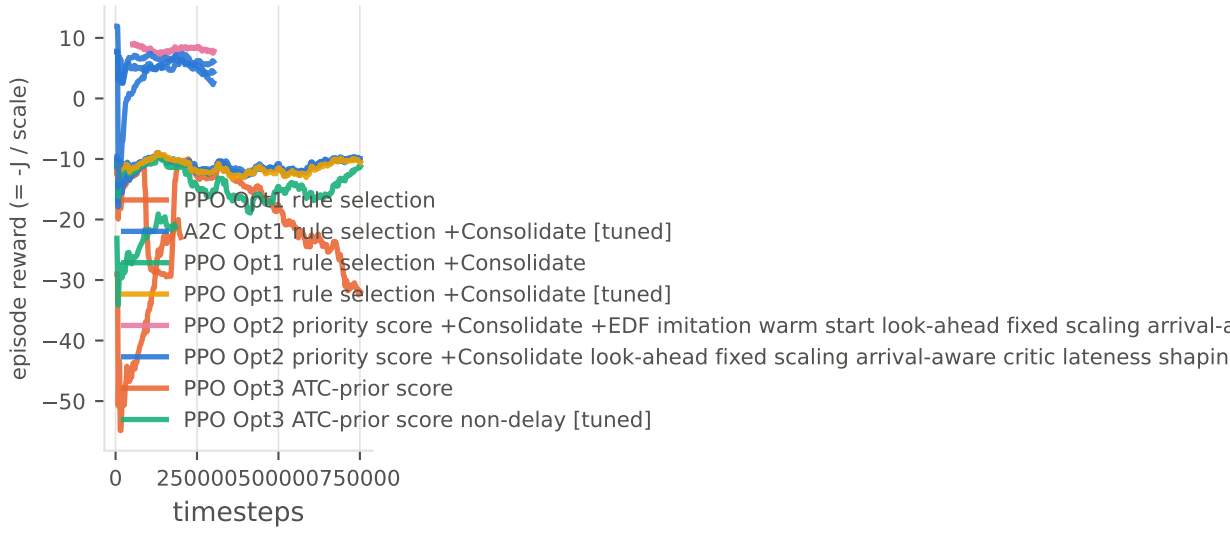
on_rho075: training curves (each line one seed)
episode reward (= -J / scale)
10
0
−10
−20
−30
−40
−50
0
250000500000750000
timesteps
PPO Opt1 rule selection
A2C Opt1 rule selection +Consolidate [tuned]
PPO Opt1 rule selection +Consolidate
PPO Opt1 rule selection +Consolidate [tuned]
PPO Opt2 priority score +Consolidate +EDF imitation warm start look-ahead fixed scaling arrival-a
PPO Opt2 priority score +Consolidate look-ahead fixed scaling arrival-aware critic lateness shapin
PPO Opt3 ATC-prior score
PPO Opt3 ATC-prior score non-delay [tuned]