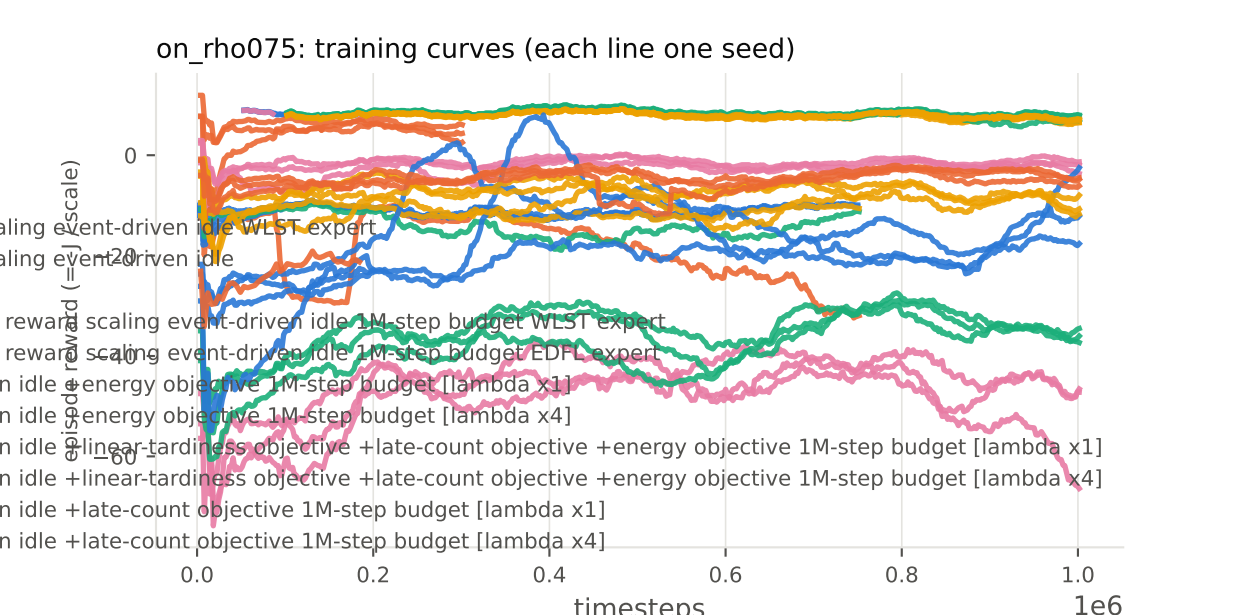

on_rho075: training curves (each line one seed)
episode reward (=J, rescale)
0
−20
−40
−60
aling event-driven idle WLST expert
aling event-driven idle
reward scaling event-driven idle 1M-step budget WLST expert
reward scaling event-driven idle 1M-step budget EDFL expert
n idle +energy objective 1M-step budget [lambda x1]
n idle +energy objective 1M-step budget [lambda x4]
n idle +linear-tardiness objective +late-count objective +energy objective 1M-step budget [lambda x1]
n idle +linear-tardiness objective +late-count objective +energy objective 1M-step budget [lambda x4]
n idle +late-count objective 1M-step budget [lambda x1]
n idle +late-count objective 1M-step budget [lambda x4]
0.0
0.2
0.4
0.6
0.8
1.0
timesteps
1e6