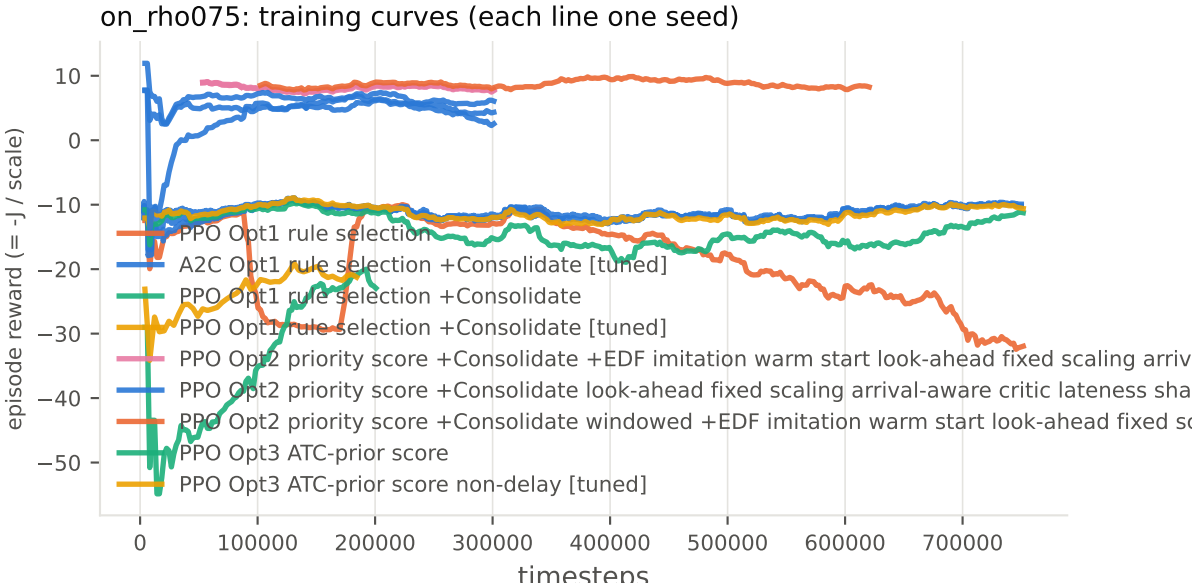
on_rho075: training curves (each line one seed)
episode reward (= -J / scale)
10
0
−10
−20
−30
−40
−50
0
100000
200000
300000
400000
500000
600000
700000
timesteps
PPO Opt1 rule selection
A2C Opt1 rule selection +Consolidate [tuned]
PPO Opt1 rule selection +Consolidate
PPO Opt1 rule selection +Consolidate [tuned]
PPO Opt2 priority score +Consolidate +EDF imitation warm start look-ahead fixed scaling arriv
PPO Opt2 priority score +Consolidate look-ahead fixed scaling arrival-aware critic lateness sha
PPO Opt2 priority score +Consolidate windowed +EDF imitation warm start look-ahead fixed sc
PPO Opt3 ATC-prior score
PPO Opt3 ATC-prior score non-delay [tuned]