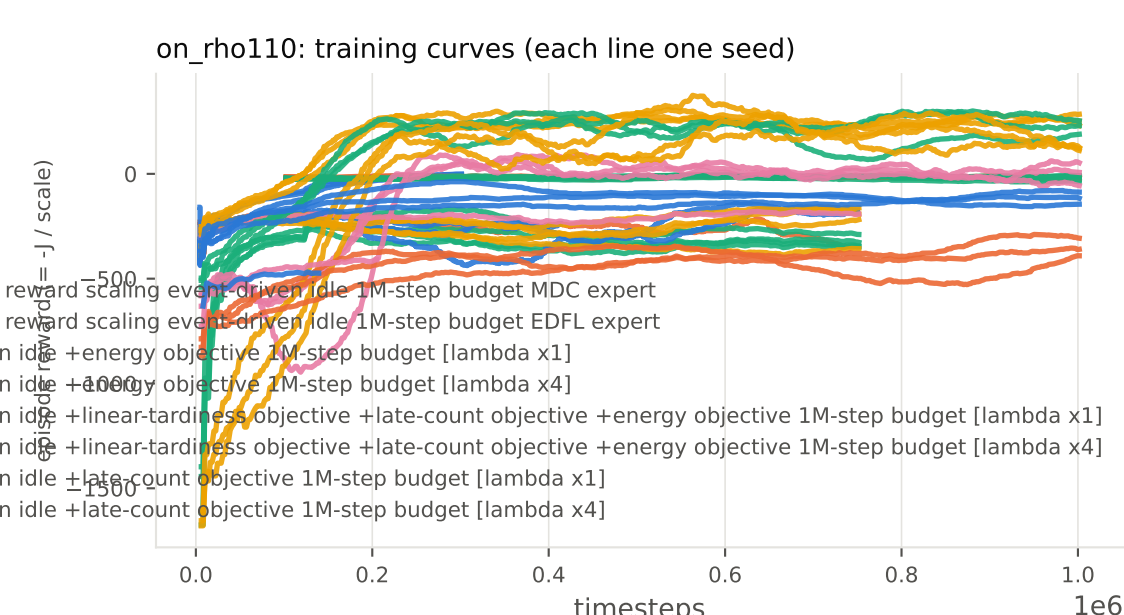

on_rho110: training curves (each line one seed)
episode reward (= -J / scale)
0
−500
−1000
−1500
reward scaling event-driven idle 1M-step budget MDC expert
reward scaling event-driven idle 1M-step budget EDFL expert
n idle +energy objective 1M-step budget [lambda x1]
n idle +energy objective 1M-step budget [lambda x4]
n idle +linear-tardiness objective +late-count objective +energy objective 1M-step budget [lambda x1]
n idle +linear-tardiness objective +late-count objective +energy objective 1M-step budget [lambda x4]
n idle +late-count objective 1M-step budget [lambda x1]
n idle +late-count objective 1M-step budget [lambda x4]
0.0
0.2
0.4
0.6
0.8
1.0
timesteps
1e6