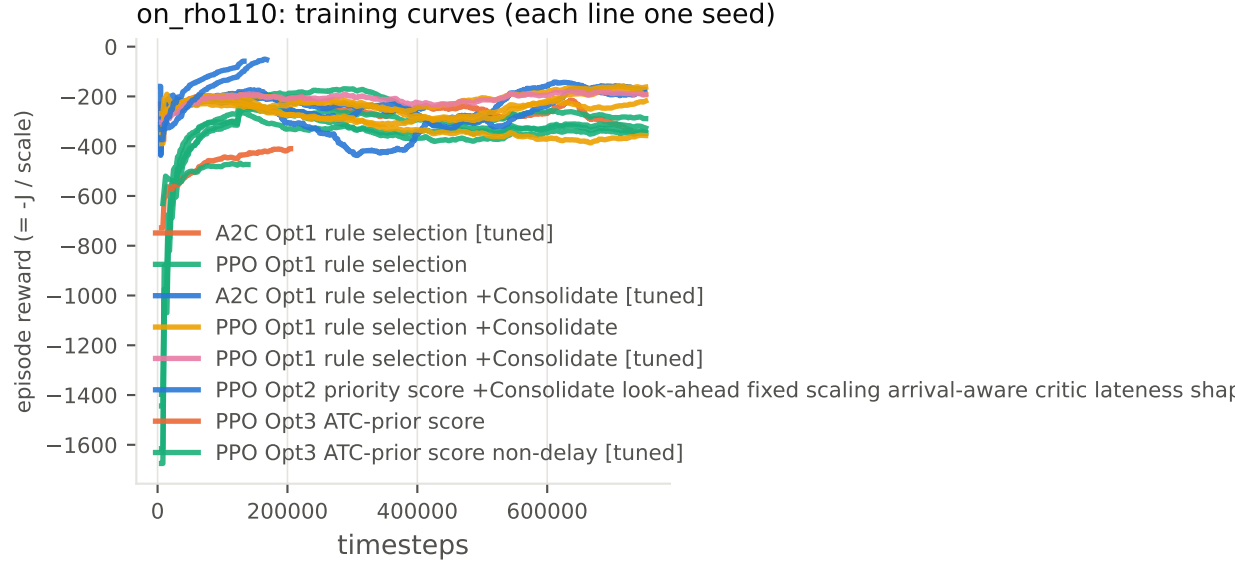
on_rho110: training curves (each line one seed)
episode reward (= -J / scale)
0
−200
−400
−600
−800
−1000
−1200
−1400
−1600
0
200000
400000
600000
timesteps
A2C Opt1 rule selection [tuned]
PPO Opt1 rule selection
A2C Opt1 rule selection +Consolidate [tuned]
PPO Opt1 rule selection +Consolidate
PPO Opt1 rule selection +Consolidate [tuned]
PPO Opt2 priority score +Consolidate look-ahead fixed scaling arrival-aware critic lateness shap
PPO Opt3 ATC-prior score
PPO Opt3 ATC-prior score non-delay [tuned]