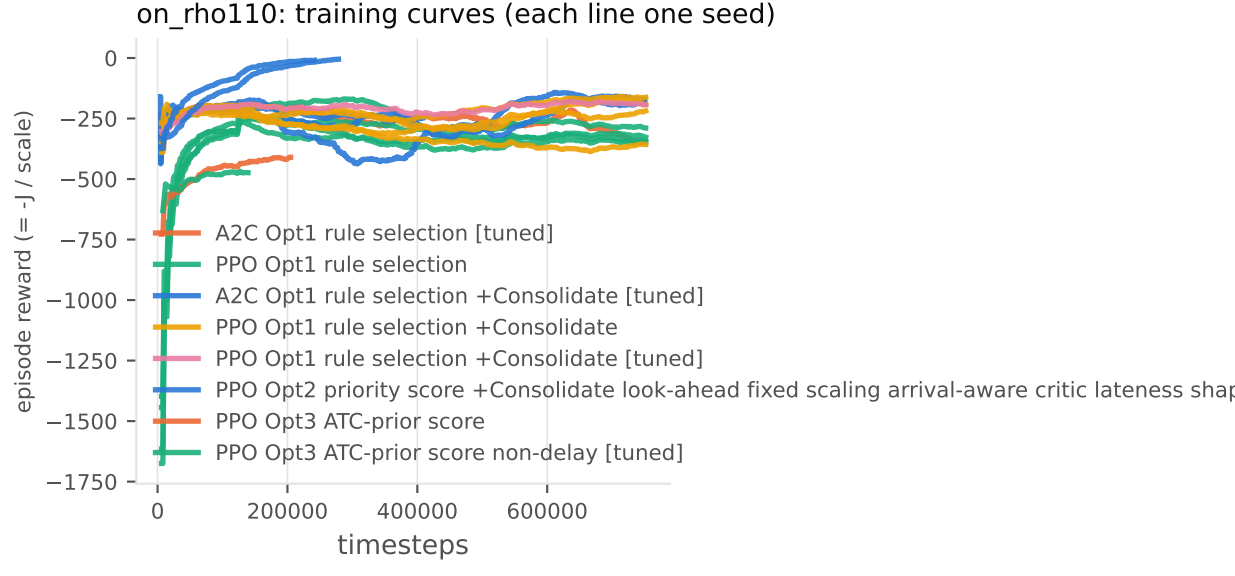
on_rho110: training curves (each line one seed)
episode reward (= -J / scale)
0
−250
−500
−750
−1000
−1250
−1500
−1750
0
200000
400000
600000
timesteps
A2C Opt1 rule selection [tuned]
PPO Opt1 rule selection
A2C Opt1 rule selection +Consolidate [tuned]
PPO Opt1 rule selection +Consolidate
PPO Opt1 rule selection +Consolidate [tuned]
PPO Opt2 priority score +Consolidate look-ahead fixed scaling arrival-aware critic lateness shap
PPO Opt3 ATC-prior score
PPO Opt3 ATC-prior score non-delay [tuned]